Miguel Fiandor Gutiérrez

AI & Data Engineering Consultant

miguel.fiandor.gutierrez@gmail.com · github.com/miguelfg · linkedin.com/in/mfiandor · miguelfg.github.io/my-portfolio

PROFILE

I design and build data and AI systems that turn messy sources into decisions. That covers scraping and extraction pipelines, graph databases, analytical dashboards and, more recently, agentic tools that run on hosted or local LLMs. Over 11 years at the ICIJ I engineered the data behind nine global investigations that won more than 50 journalism awards. I now help teams design, build and integrate GenAI solutions that hold up in production.

SERVICES

01 **GenAI solutions & agentic systems.** Design, develop and integrate GenAI into products and workflows, from feasibility study to production, on hosted or on-premise local models.

02 **Data pipelines & extraction.** Web scraping, PDF parsing, standardization and ETL from millions of documents into structured stores.

03 **Graph & investigative data.** Entity resolution, Neo4j modeling and OSINT research for fraud, compliance and journalism.

04 **Dashboards & internal tools.** Streamlit and Metabase analytics, CLIs and AI skills that non-engineers can operate.

SELECTED ENGAGEMENTS

Ethon Shield · Cybersecurity

AI Lead, GenAI feasibility study · Jun 2026 → Oct 2026

Assessed whether GenAI can support baseband firmware security analysis. I defined the scope and success criteria, designed the analysis workflow, and built and evaluated a proof of concept on local models (llama.cpp, Qwen).

International Consortium of Investigative Journalists (ICIJ)

Data Analyst & Engineer, Tech and Data & Research teams · Jun 2015 → Oct 2026

Engineered data infrastructure for nine major cross-border investigations, including the Panama Papers, the Pandora Papers and the Coin Laundry crypto investigation. I converted millions of leaked documents into structured and graph databases, built scraping and PDF-extraction pipelines and internal dashboards, and improved research and search workflows. I also contributed to open-source tools Prophecies and Datashare Tarentula. **Outcome:** data foundations for work recognized with 50+ awards, including a Pulitzer Prize.

Python · Pandas · Neo4j · Linkurious · Streamlit · AWS · NLP · Docker

AVAILABILITY

Remote · EU and US East time zones
Full or part time, or project basis

CORE STACK

Language & data
Python · Pandas · NumPy · Jupyter · REST / OpenAPI

GenAI & agents
Claude · Codex · MCP · CrewAI · LangChain · AI skills

Local models
llama.cpp · Ollama · Qwen

ML & NLP
Scikit-Learn · spaCy

Storage & graph
Neo4j · PostgreSQL · SQLite · Elasticsearch

Apps & BI
Streamlit · Django · Metabase · Linkurious

Cloud & ops
GCP Cloud Run · AWS · Docker · Git

CERTIFICATIONS

Neo4j Certified Professional 2018
Elastic Certified Engineer I 2020
OSINT Judicial Expert, UNED 2021

LANGUAGES

Spanish, native
English, professional

SELECTED PROJECTS

CourtListener CLI & AI skill

A Python CLI for the CourtListener legal API with 18 command groups, batch processing and JSON/CSV/XLSX export. Its companion skill for Claude, Codex and Gemini turns plain-English legal research requests into CLI actions, which cuts inference-token use by design.

API-to-CLI skillset

Three published AI skills that turn API documentation into an OpenAPI spec, a PRD and a Python CLI scaffold, built for analysts and researchers working in Claude Code or Codex.

Agentic Text2SQL

A CrewAI and Streamlit app that lets non-technical users query business databases in natural language. A three-agent pipeline raises SQL quality, and a batch CLI benchmarks models and providers across databases.

Spanish public property auctions

A pipeline that collects and normalizes BOE auction data, with an analytical dashboard on Google Cloud Run and collection jobs run by Cloud Run Jobs and Scheduler.

Cancer Calculus

Clinical-trial research for a published project, using a ClinicalTrials.gov MCP server and Claude's extended research.

Public-data dashboards

Deployed analytics for Toyota used-car sales in Spain, Roadsurfer European spots and Madrid daily pollen.

EARLIER CAREER

2014 → 2015 **Freelance Python/Django developer.** Built web platforms and REST APIs for two startups, Apparcando and Estigi.

2013 → 2015 **Founder, Transparencia de Cuentas Públicas.** Ran an open-data platform on Spanish public accounts.

2010 → 2012 **IT manager, Acciona Energy.** Ran trading-system infrastructure, an internal BI rollout and reporting automation.

EDUCATION

2001 → 2009 **MSc Computer Engineering**, Universidad Politécnica de Madrid. Exchange year at the University of Copenhagen (DIKU).

2009 → 2010 **MSc Renewable Energies & Environment**, Universidad Politécnica de Madrid (ETSIDI).

2022 → 2024 **Continuing education.** DataCamp Machine Learning Scientist track (25 courses) and DeepLearning.ai courses on LLM apps, agents and LangGraph.

TALKS & MENTIONS

GIJN Cited in the Global Investigative Journalism Network's Reporter's Guide to Investigating Cryptocurrency, The People and the Technology Behind the Panama Papers and How a Global Team Used Data to Interrogate Green Claims for Deforestation Inc.

2024 Graphs4Good: Building a better world with connected data, Neo4j Connections

2021 Las 4 Ps: Python y Pandas en Pandora Papers

2017 El BackOffice de un Premio Pulitzer: Los Papeles de Panamá, Neo4j Masterclass Madrid